

Announcements

- ❖ **Project 2: Multi-agent Search**
 - ❖ Due Sunday June 20 at 11:59pm.
- ❖ **Homework 4: RL**
 - ❖ Has been released! Due Wed. June 23 at 11:59pm.
- ❖ **Midterm Exam**
 - ❖ June 25, 12:10pm-1:50pm
 - ❖ Covers everything, CSP included.
- ❖ **Recitations**
 - ❖ June 22, 1pm-3pm; followed by OH

Ve492: Introduction to Artificial Intelligence

Constraint Satisfaction Problems I



Paul Weng

UM-SJTU Joint Institute

Slides adapted from <http://ai.berkeley.edu>, AIMA, UM, CMU

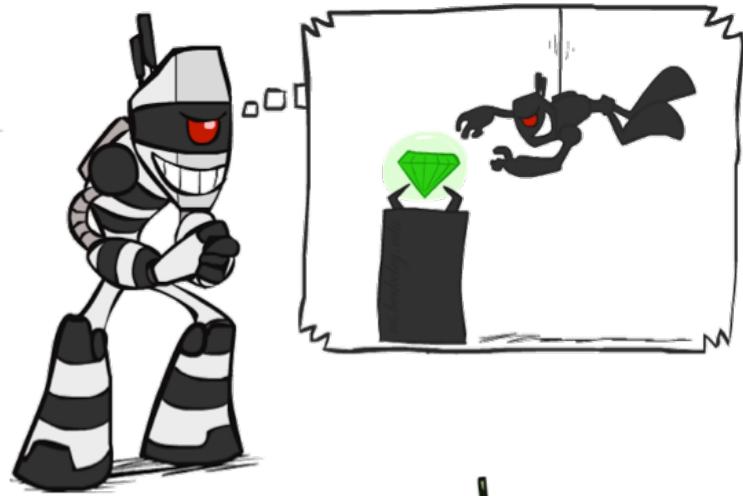
What have we learned so far?

- ❖ Search and planning
 - ❖ Define a state space, goal test; Find path from start to goal
- ❖ Game trees
 - ❖ Define utilities; Find path from start that maximizes utility
- ❖ Decision theory and game theory
 - ❖ Foundation for MEU; Basic concepts in game theory
- ❖ MDPs
 - ❖ Define rewards, utility = (discounted) sum of rewards
 - ❖ Find policy that maximizes utility
- ❖ Reinforcement learning
 - ❖ Just like MDPs, only T and / or R are not known in advance
- ❖ Today: constraint satisfaction
 - ❖ Find solution that satisfies constraints; Not just for finding a sequential plan

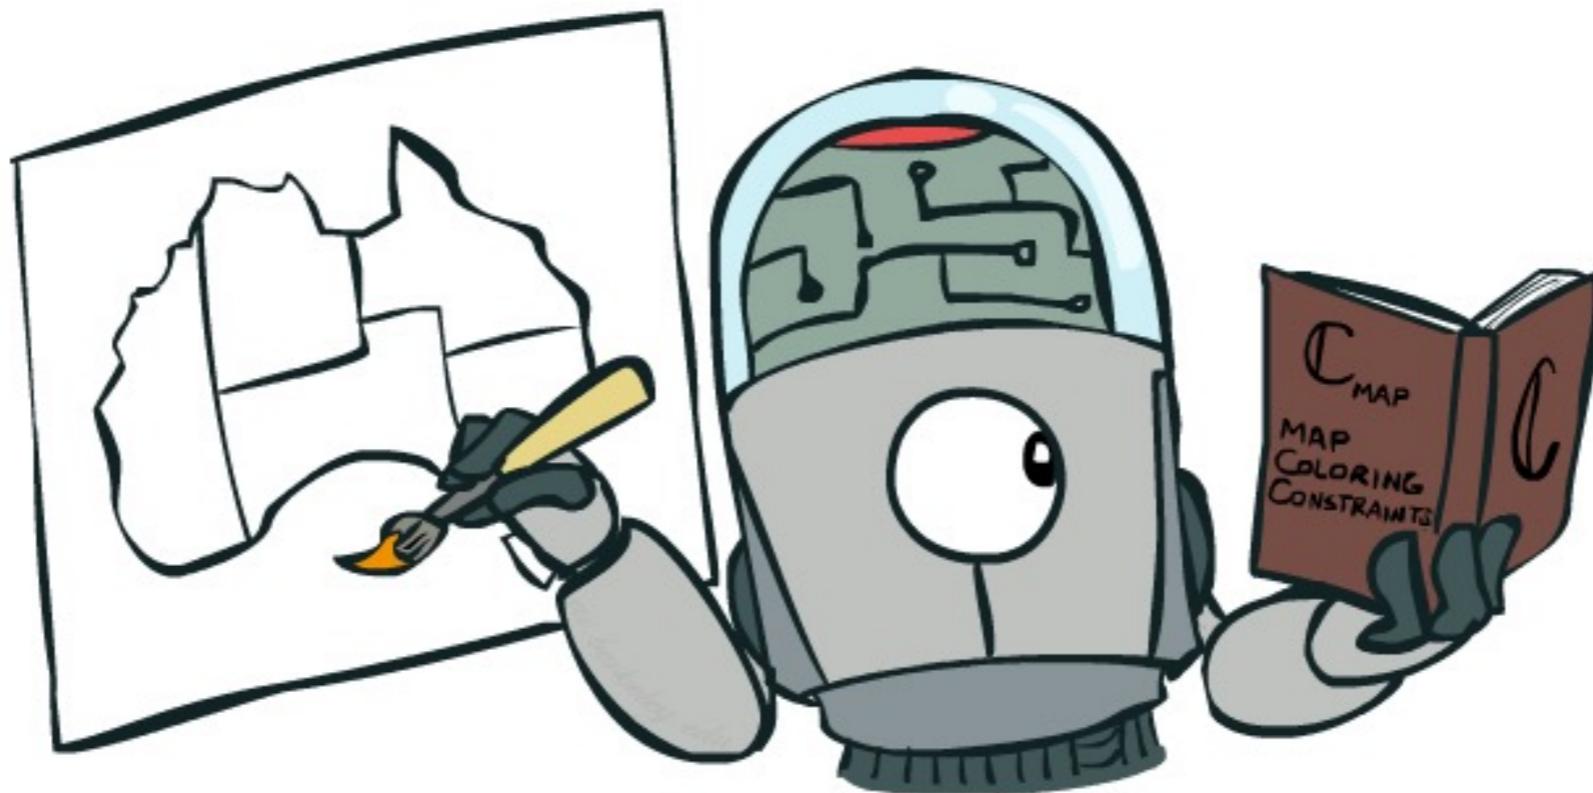


What is Search For?

- ❖ Assumptions about the world: a single agent, deterministic actions, fully observed state, discrete state space
- ❖ Planning: sequences of actions
 - ❖ The path to the goal is the important thing
 - ❖ Paths have various costs, depths
 - ❖ Heuristics give problem-specific guidance
- ❖ Identification: assignments to variables
 - ❖ The goal itself is important, not the path
 - ❖ All paths at the same depth (for some formulations)
 - ❖ CSPs are specialized for identification problems

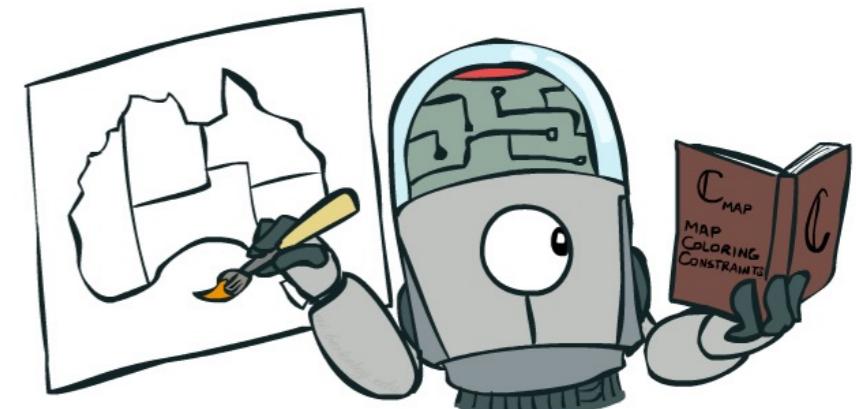
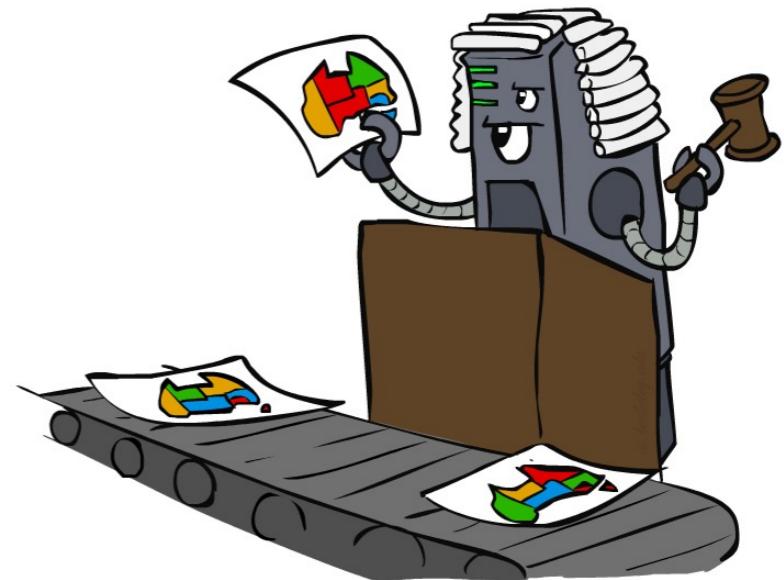


Constraint Satisfaction Problems

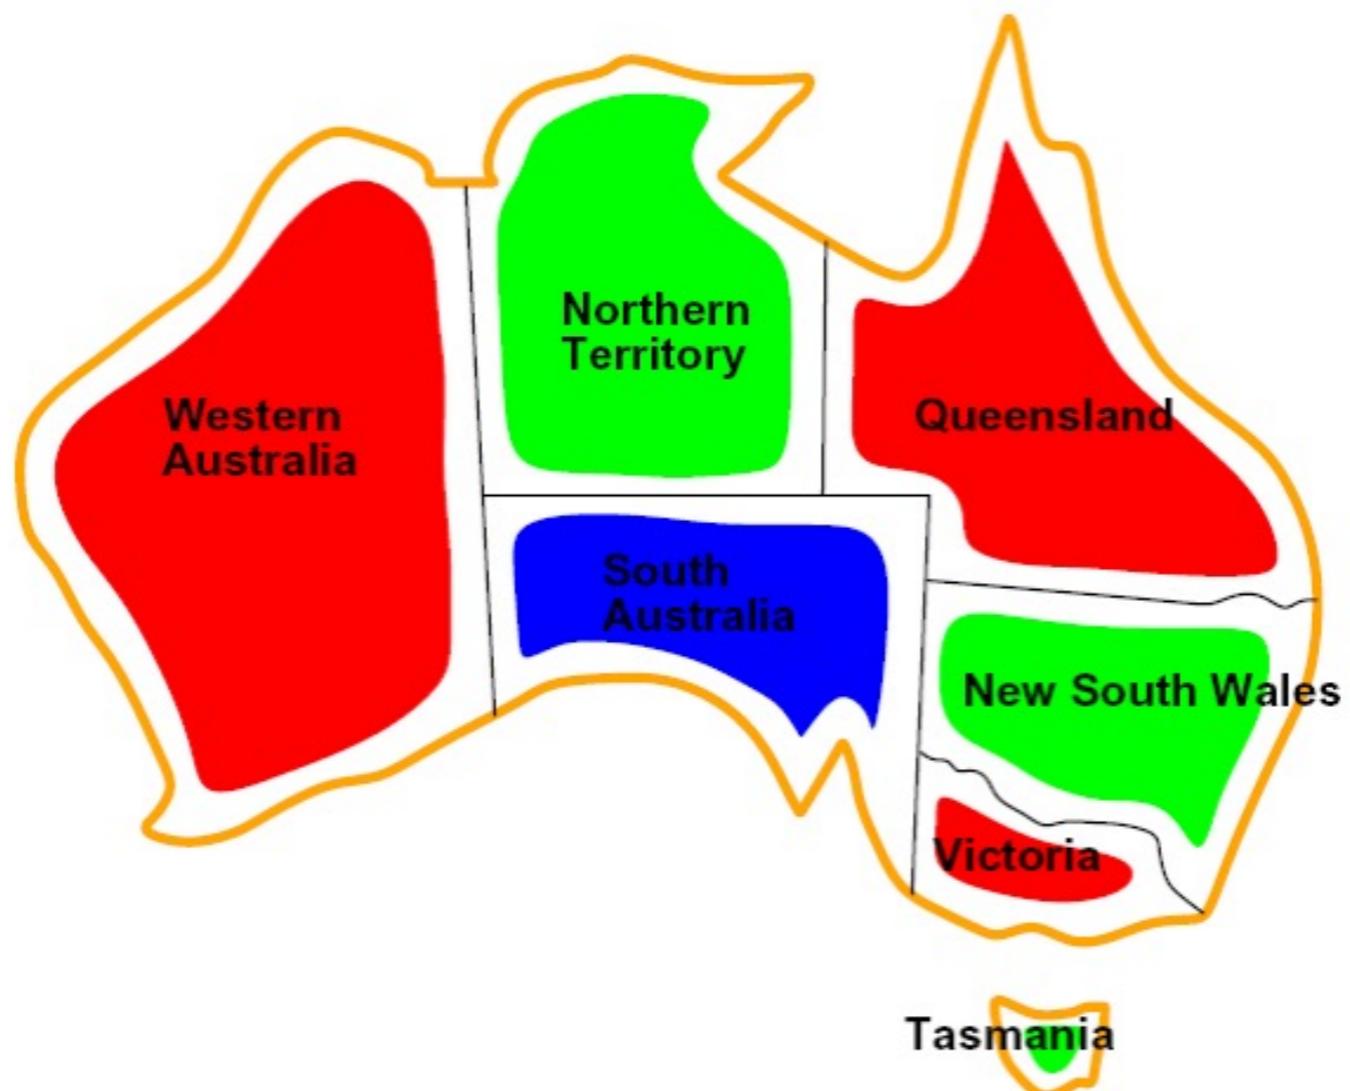


Constraint Satisfaction Problems

- ❖ Standard search problems:
 - ❖ State is a “black box”: arbitrary data structure
 - ❖ Goal test can be any function over states
 - ❖ Successor function can also be anything
- ❖ Constraint satisfaction problems (CSPs):
 - ❖ A special subset of search problems
 - ❖ State is defined by variables X_i with values from a domain D (sometimes D depends on i)
 - ❖ Goal test is a set of constraints specifying allowable combinations of values for subsets of variables
- ❖ Simple example of a *formal representation language*
- ❖ Allows useful general-purpose algorithms with more power than standard search algorithms



CSP Examples



Example: Map Coloring

- ❖ Variables: WA, NT, Q, NSW, V, SA, T

- ❖ Domains: $D = \{\text{red, green, blue}\}$

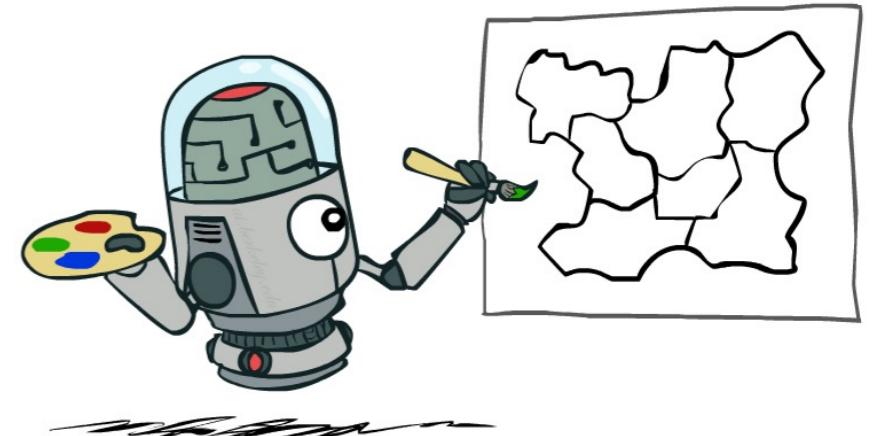
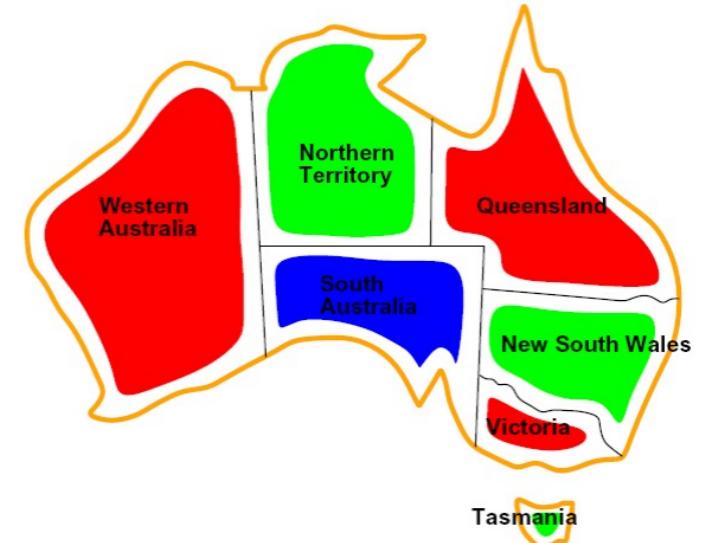
- ❖ Constraints: adjacent regions must have different colors

Implicit: $\text{WA} \neq \text{NT}$

Explicit: $(\text{WA}, \text{NT}) \in \{(\text{red, green}), (\text{red, blue}), \dots\}$

- ❖ Solutions are assignments satisfying all constraints, e.g.:

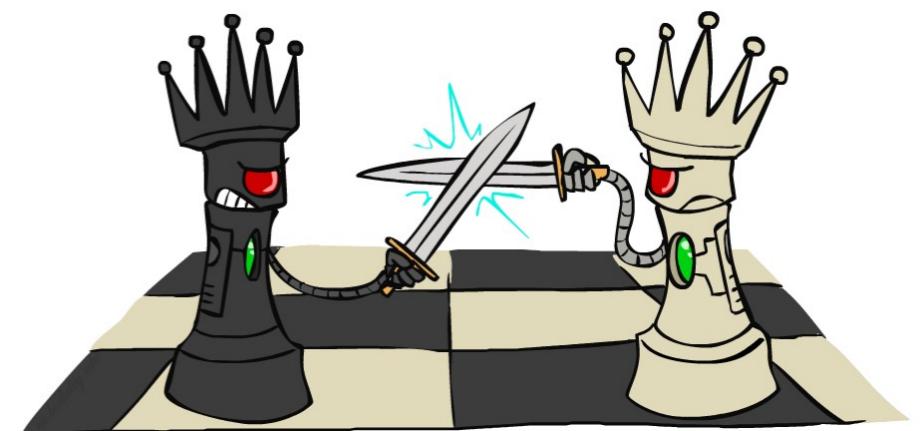
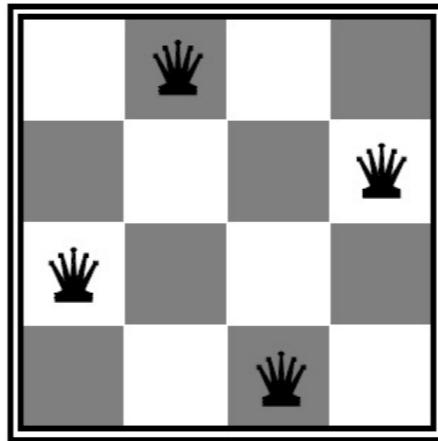
$\{\text{WA=red, NT=green, Q=red, NSW=green, V=red, SA=blue, T=green}\}$



Example: N-Queens

❖ Formulation 1:

- ❖ Variables: X_{ij}
- ❖ Domains: $\{0, 1\}$
- ❖ Constraints



$$\forall i, j, k \quad (X_{ij}, X_{ik}) \in \{(0,0), (0,1), (1,0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{kj}) \in \{(0,0), (0,1), (1,0)\}$$

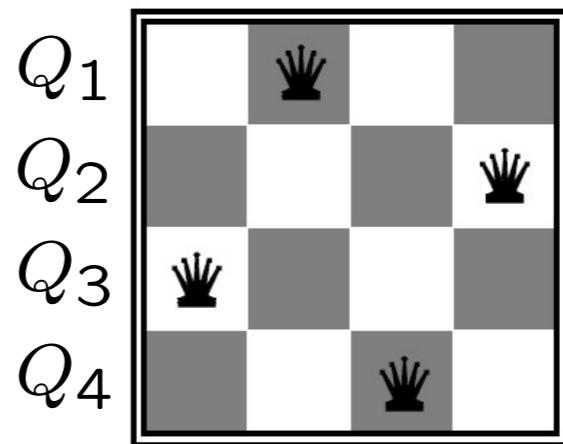
$$\forall i, j, k \quad (X_{ij}, X_{i+k,j+k}) \in \{(0,0), (0,1), (1,0)\}$$

$$\forall i, j, k \quad (X_{ij}, X_{i+k,j-k}) \in \{(0,0), (0,1), (1,0)\}$$

$$\sum_{i,j} X_{ij} = N$$

Example: N-Queens

- ❖ Formulation 2:
 - ❖ Variables: Q_k
 - ❖ Domains: $\{1, 2, 3, \dots, N\}$
 - ❖ Constraints:

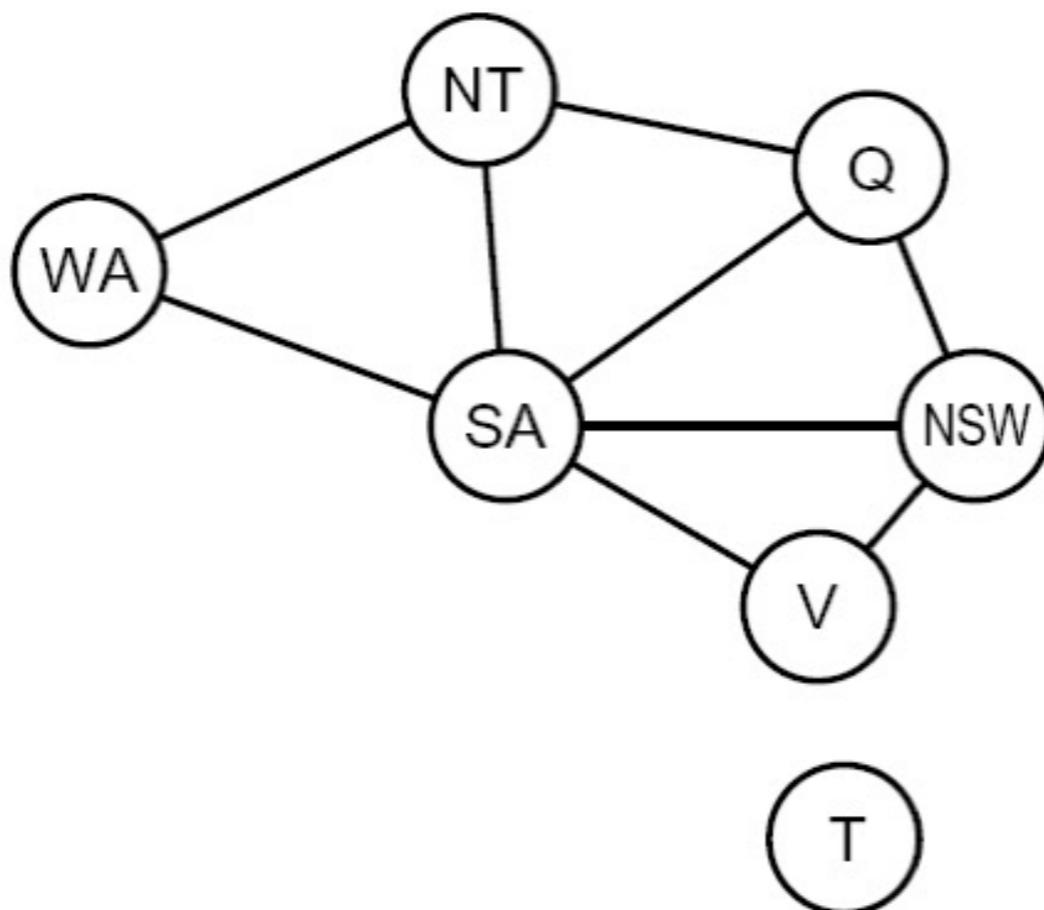


Implicit: $\forall i, j \text{ non-threatening}(Q_i, Q_j)$

Explicit: $(Q_1, Q_2) \in \{(1, 3), (1, 4), \dots\}$

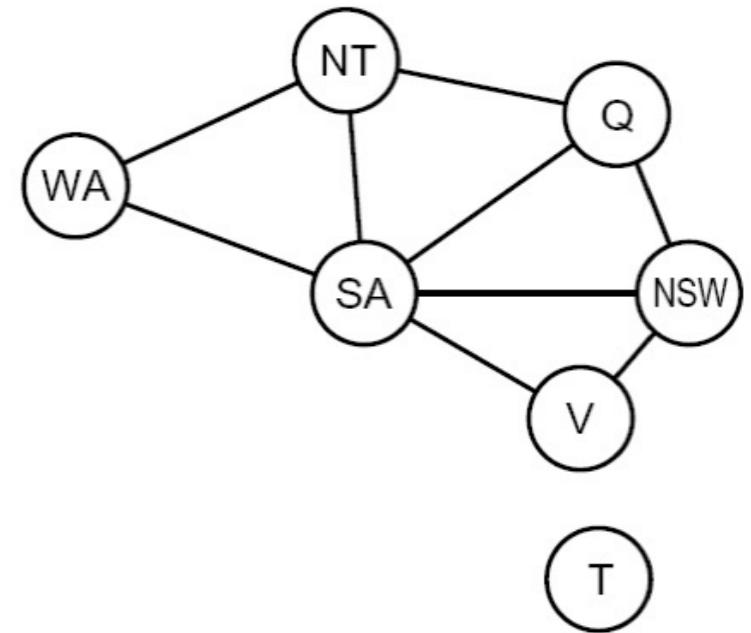
...

Constraint Graphs



Constraint Graphs

- ❖ Binary CSP: each constraint relates (at most) two variables
- ❖ Binary constraint graph: nodes are variables, arcs show constraints
- ❖ General-purpose CSP algorithms use the graph structure to speed up search. E.g., Tasmania is an independent subproblem!
- ❖ Example: [CSP solver](#)



Example: Cryptarithmetic

- ❖ **Variables:**

$F \ T \ U \ W \ R \ O \ X_1 \ X_2 \ X_3$

- ❖ **Domains:**

$\{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$

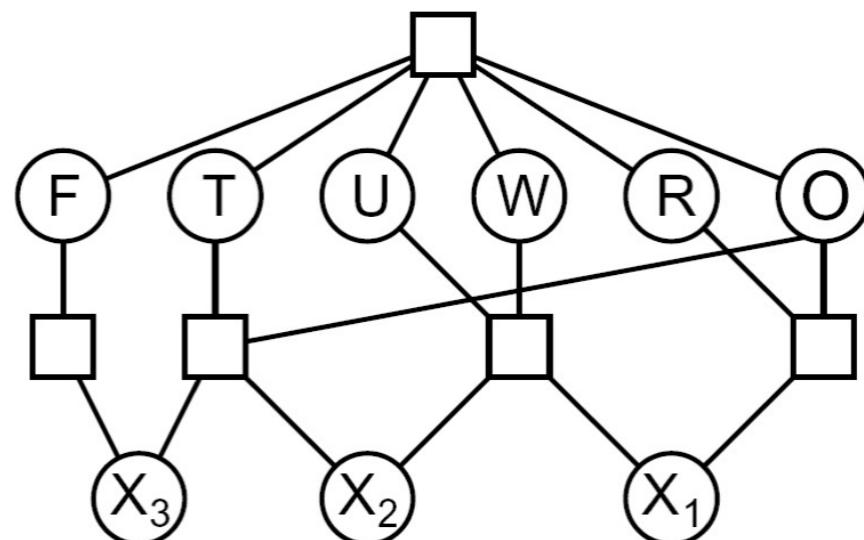
- ❖ **Constraints:**

$\text{alldiff}(F, T, U, W, R, O)$

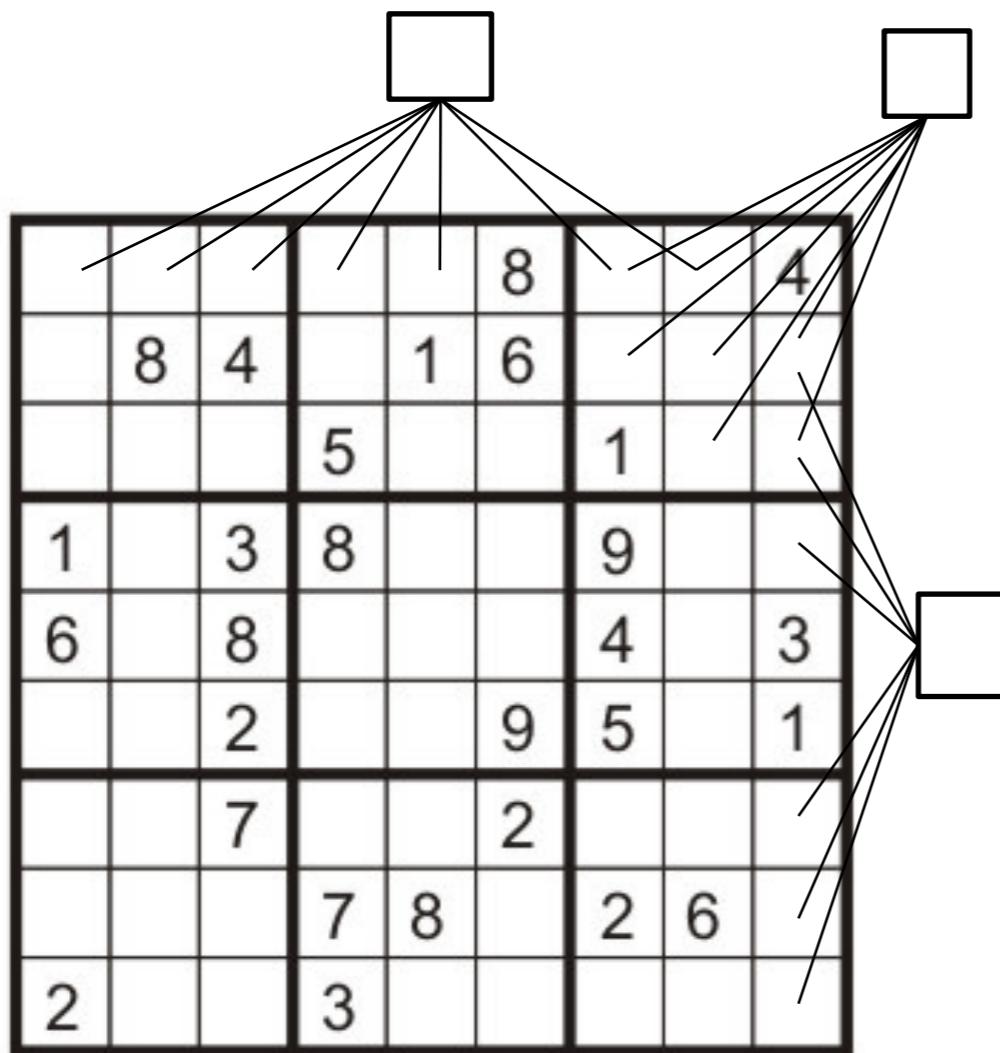
$$O + O = R + 10 \cdot X_1$$

...

$$\begin{array}{r} \text{T} \ \text{W} \ \text{O} \\ + \ \text{T} \ \text{W} \ \text{O} \\ \hline \text{F} \ \text{O} \ \text{U} \ \text{R} \end{array}$$

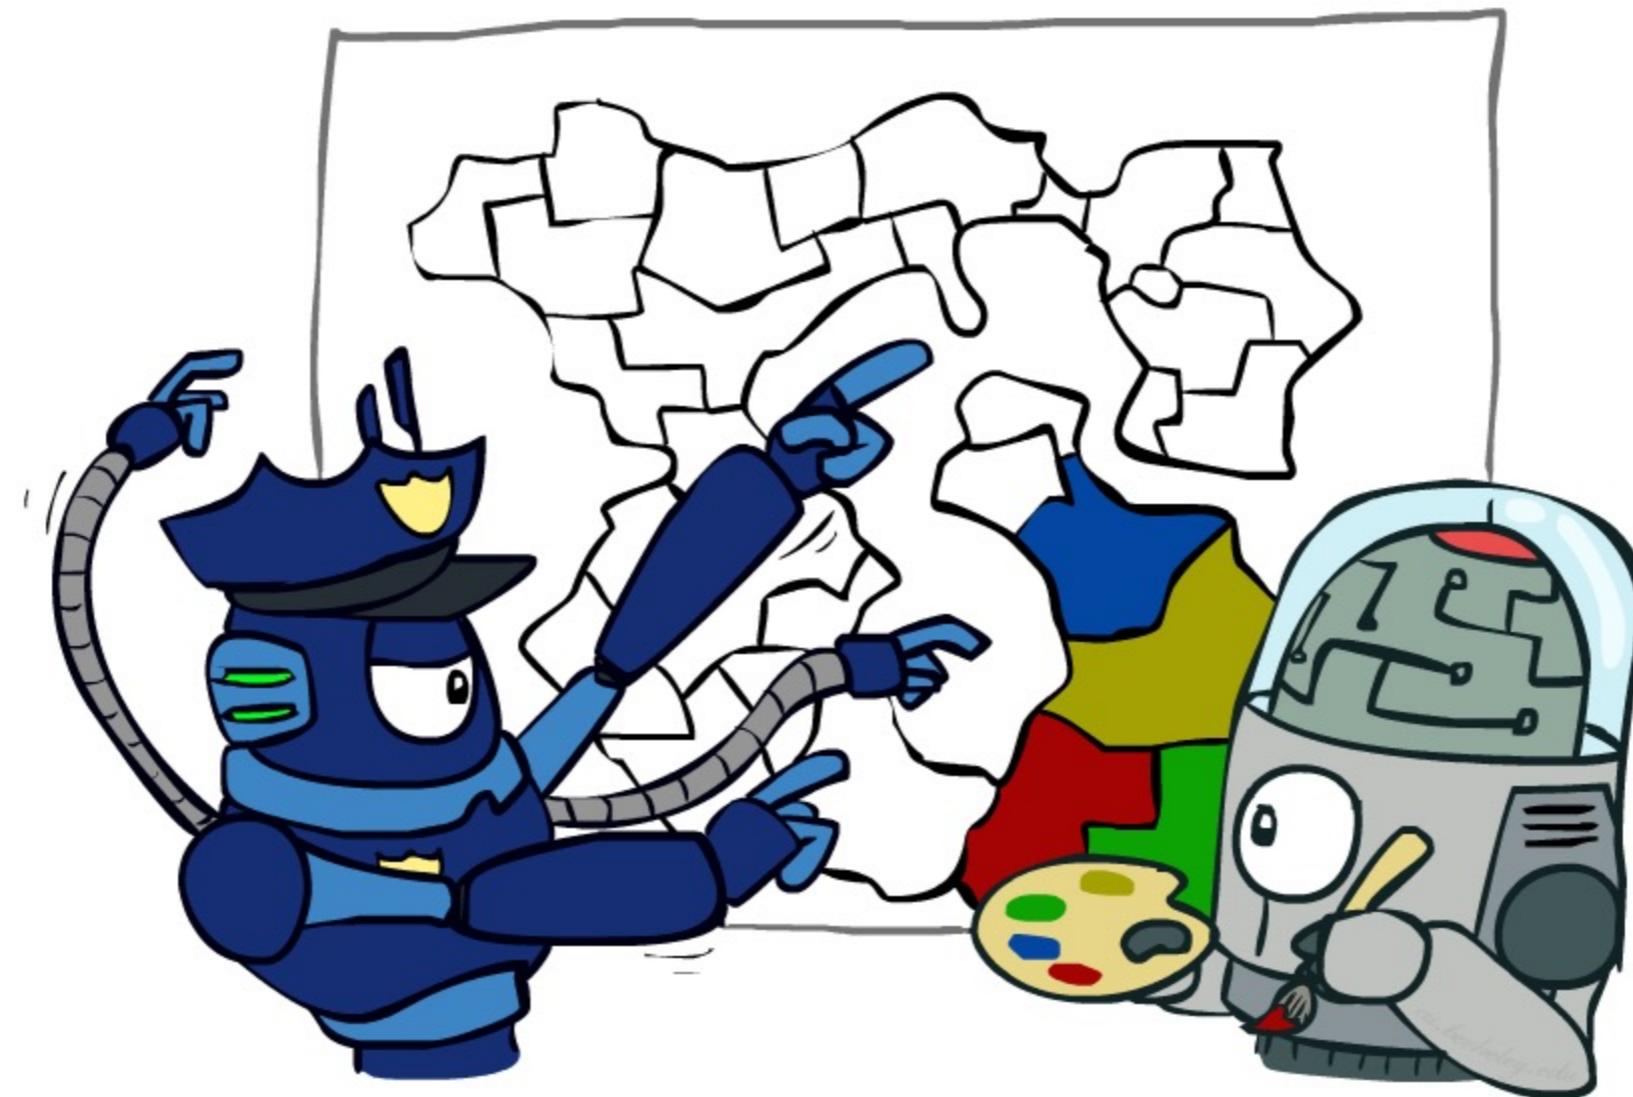


Example: Sudoku



- ❖ **Variables:**
Each (open) square
- ❖ **Domains:**
 $\{1, 2, \dots, 9\}$
- ❖ **Constraints:**
 - 9-way alldiff for each column
 - 9-way alldiff for each row
 - 9-way alldiff for each region
 - (or can have a bunch of pairwise inequality constraints)

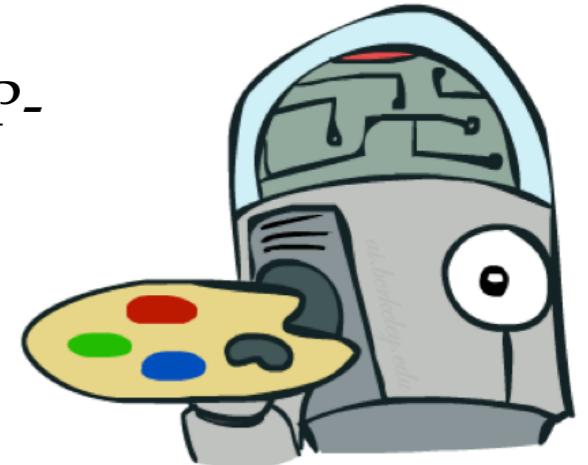
Varieties of CSPs and Constraints



Varieties of CSPs

❖ Discrete Variables

- ❖ Finite domains
 - ❖ Size d means $O(d^n)$ complete assignments
 - ❖ E.g., Boolean CSPs, including Boolean satisfiability (NP-complete)
- ❖ Infinite domains (integers, strings, etc.)
 - ❖ E.g., job scheduling, variables are start/end times for each job
 - ❖ Linear constraints solvable, nonlinear undecidable



❖ Continuous variables

- ❖ E.g., start/end times for Hubble Telescope observations
- ❖ Linear constraints solvable in polynomial time by LP methods (see Ve555 for a bit of this theory)



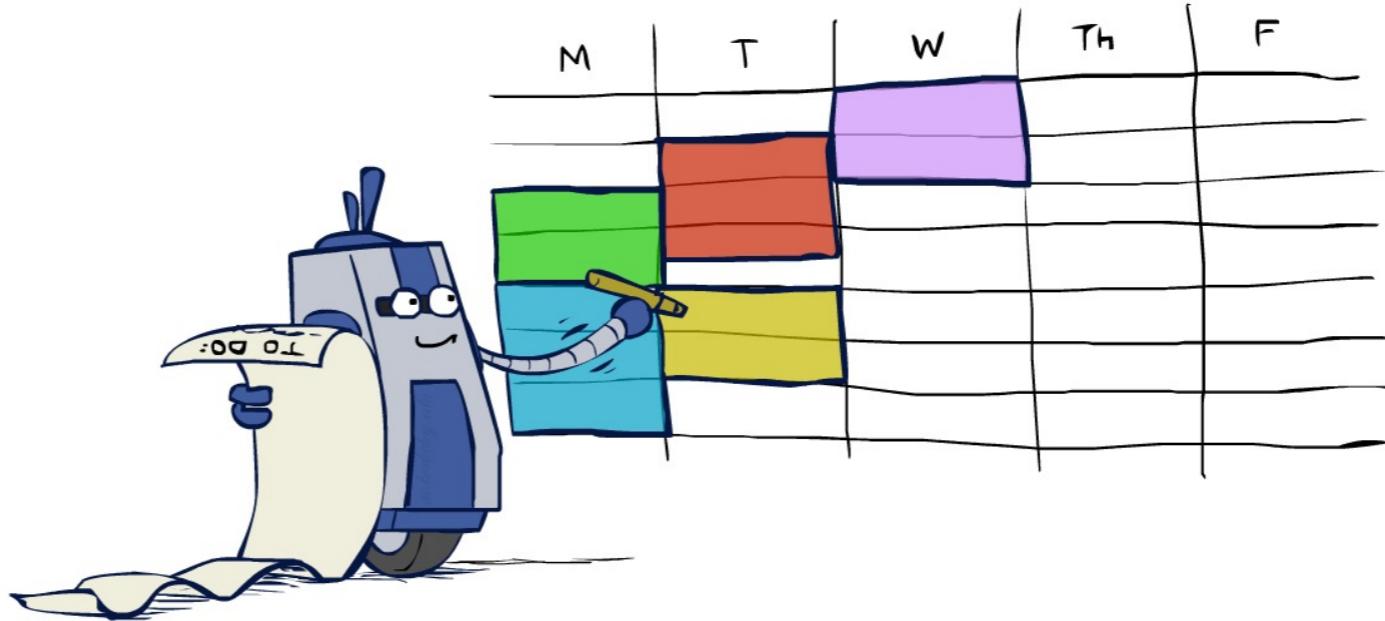
Varieties of Constraints

- ❖ Varieties of Constraints
 - ❖ Unary constraints involve a single variable (equivalent to reducing domains), e.g.:
 $SA \neq \text{green}$
 - ❖ Binary constraints involve pairs of variables, e.g.:
 $SA \neq WA$
 - ❖ Higher-order constraints involve 3 or more variables:
e.g., cryptarithmic column constraints
- ❖ Preferences (soft constraints):
 - ❖ E.g., red is better than green
 - ❖ Often representable by a cost for each variable assignment
 - ❖ Gives constrained optimization problems
 - ❖ (We'll ignore these until we get to Bayes' nets)



Real-World CSPs

- ❖ Assignment problems: e.g., who teaches what class
- ❖ Timetabling problems: e.g., which class is offered when and where?
- ❖ Hardware configuration
- ❖ Transportation scheduling
- ❖ Factory scheduling
- ❖ Circuit layout
- ❖ Fault diagnosis
- ❖ ... lots more!



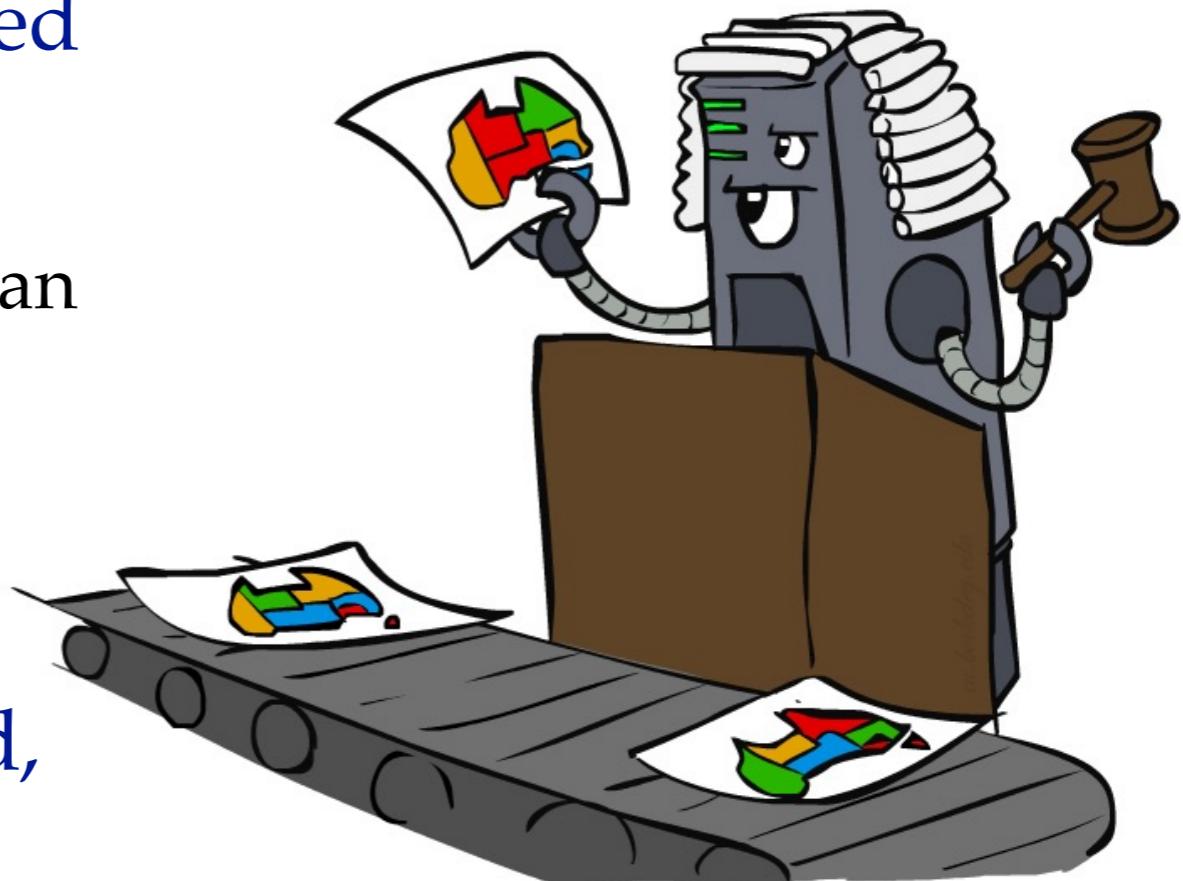
- ❖ Many real-world problems involve real-valued variables...

Solving CSPs



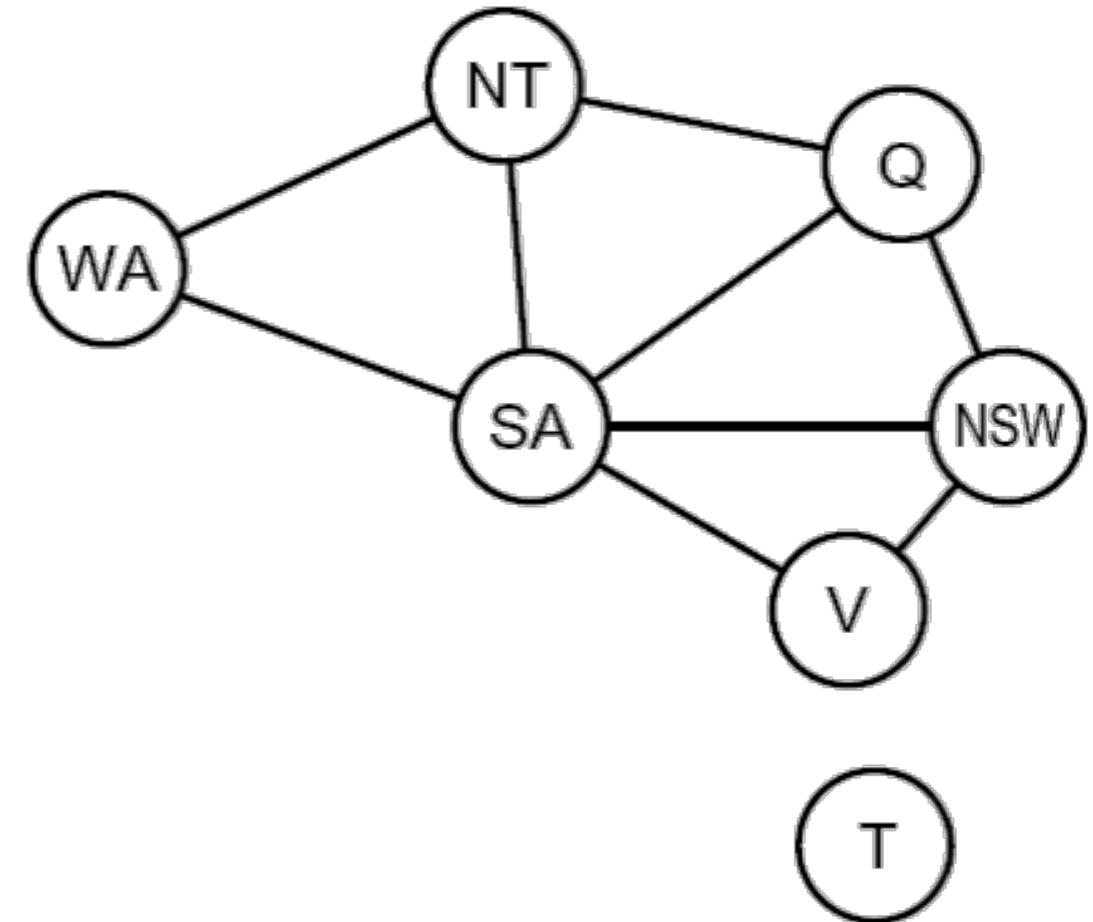
Standard Search Formulation

- ❖ Standard search formulation of CSPs
- ❖ States defined by the values assigned so far (partial assignments)
 - ❖ Initial state: the empty assignment, {}
 - ❖ Successor function: assign a value to an unassigned variable
 - ❖ Goal test: the current assignment is complete and satisfies all constraints
- ❖ We'll start with the straightforward, naïve approach, then improve it



Search Methods

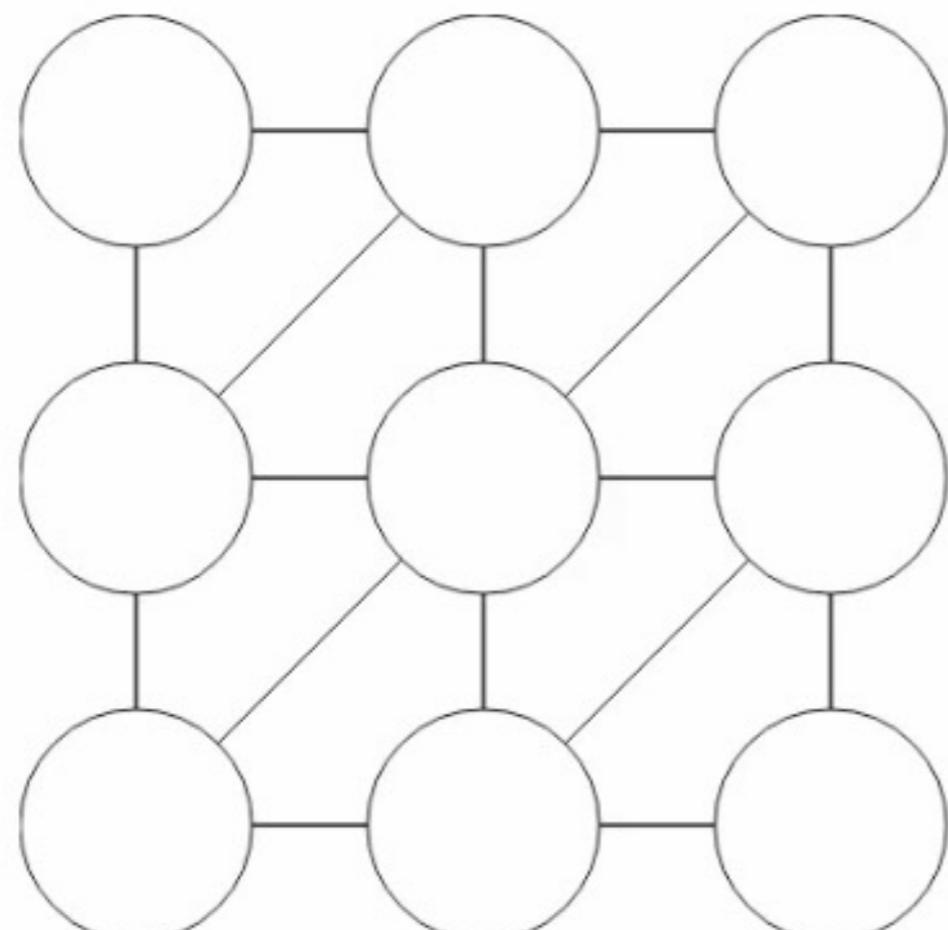
- ❖ What would BFS do?



- ❖ What would DFS do?

- ❖ What problems does naïve search have?

Video of Demo Coloring -- DFS



Reset

Prev

Pause

Next

Play

Faster

Graph

Simple

Algorithm

Naive Search

Ordering

- None
- MRV
- MRV with LCV

Filtering

- None
- Forward Checking
- Arc Consistency

Speed

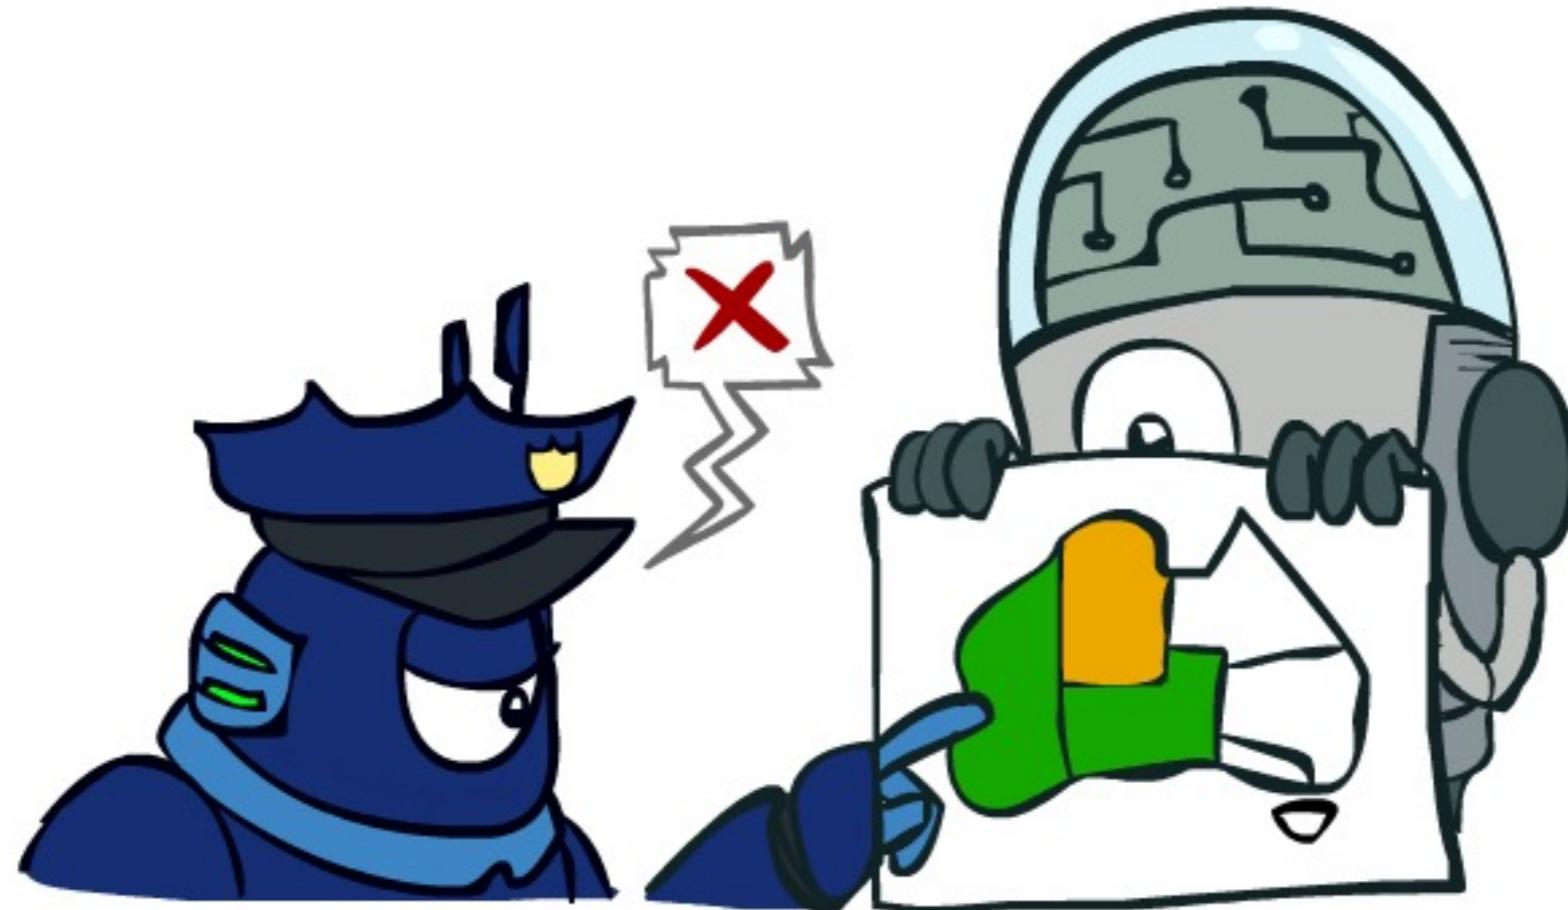
Speedup

1

Frame Delay

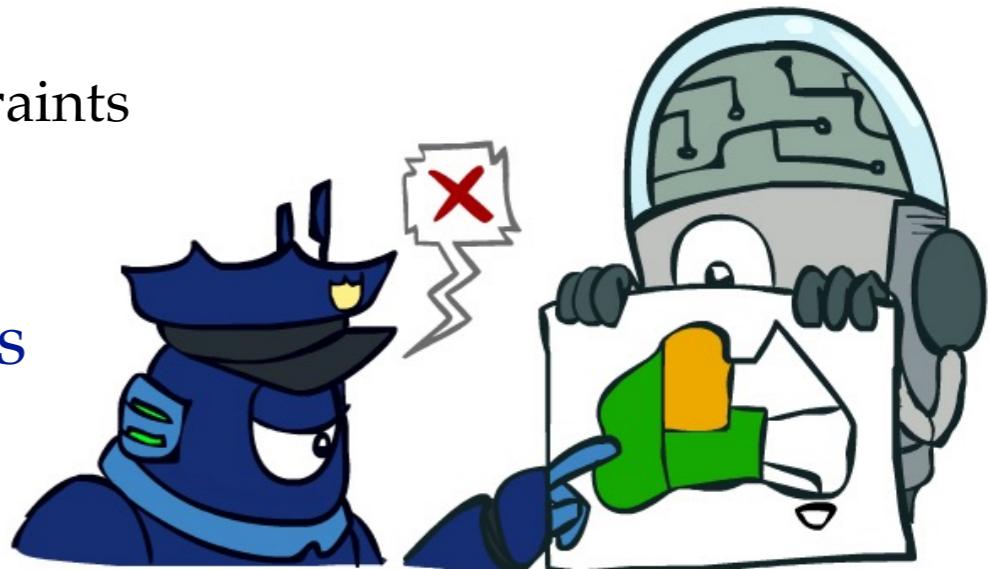
700

Backtracking Search

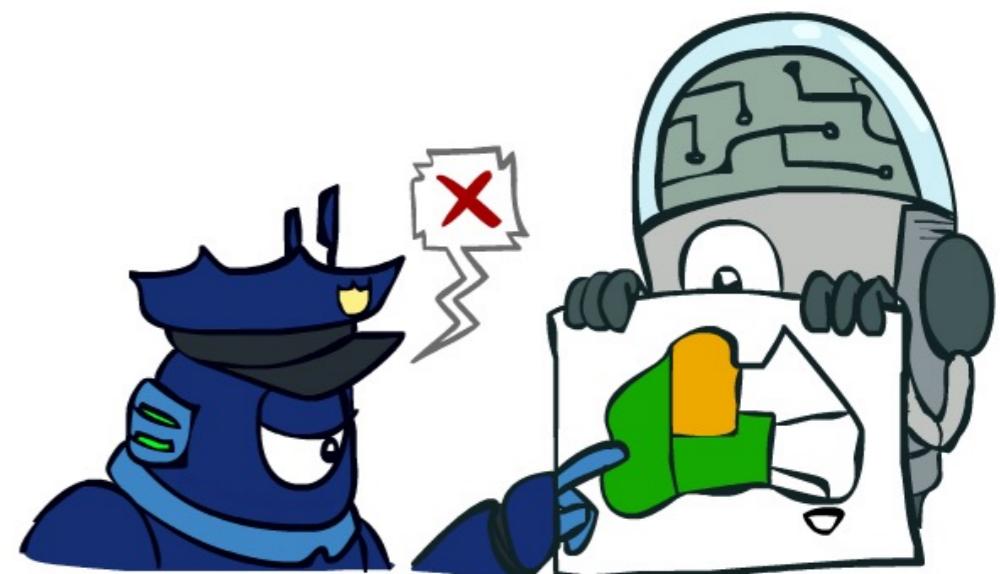
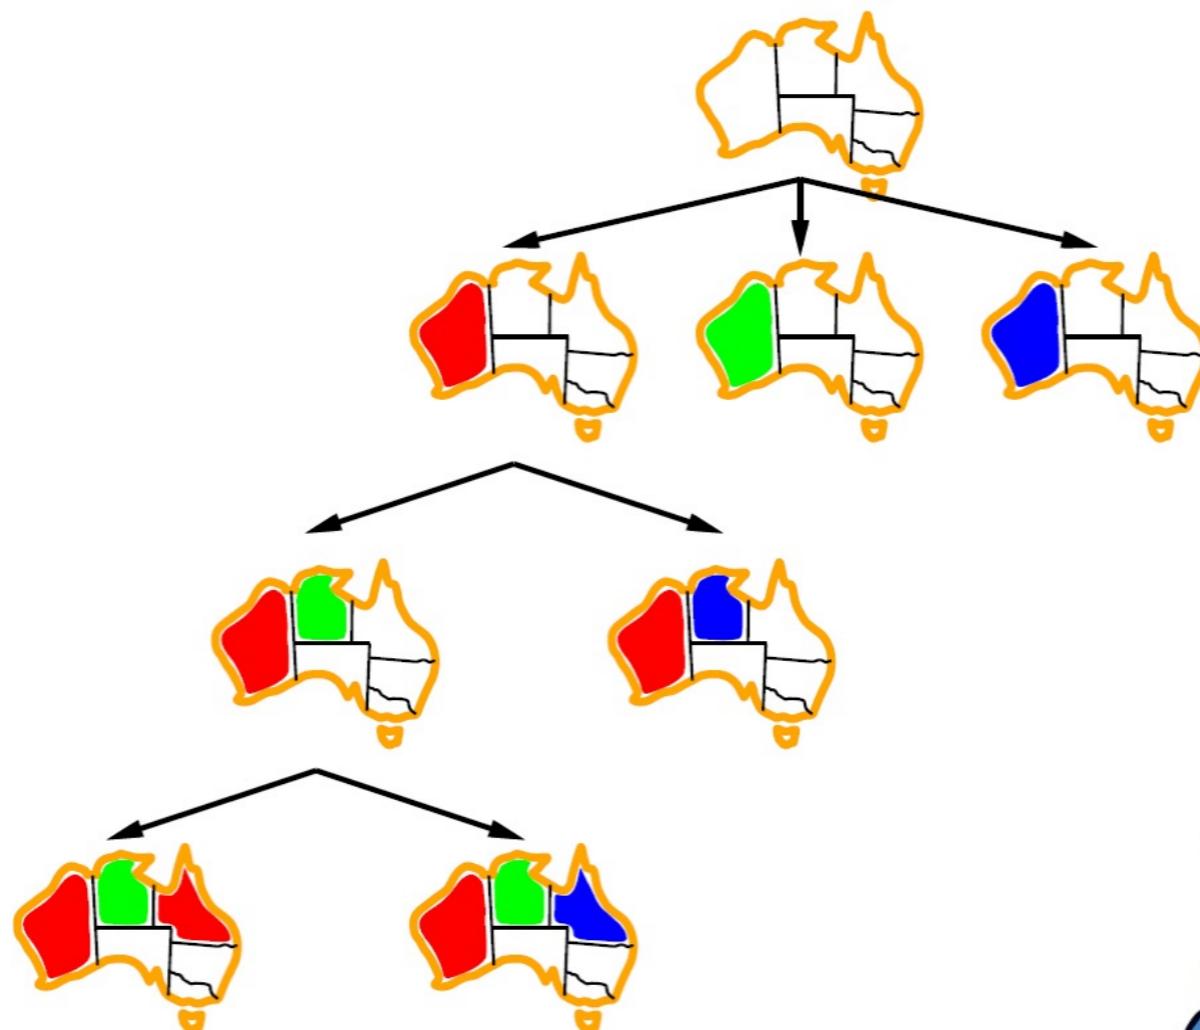


Backtracking Search

- ❖ Backtracking search is the basic uninformed algorithm for solving CSPs
- ❖ Idea 1: One variable at a time
 - ❖ Variable assignments are commutative, so fix ordering
 - ❖ I.e., [WA = red then NT = green] same as [NT = green then WA = red]
 - ❖ Only need to consider assignments to a single variable at each step
- ❖ Idea 2: Check constraints as you go
 - ❖ I.e. consider only values which do not conflict previous assignments
 - ❖ Might have to do some computation to check the constraints
 - ❖ “Incremental goal test”
- ❖ Depth-first search with these two improvements is called *backtracking search* (not the best name)
- ❖ Can solve n-queens for $n \approx 25$



Backtracking Example

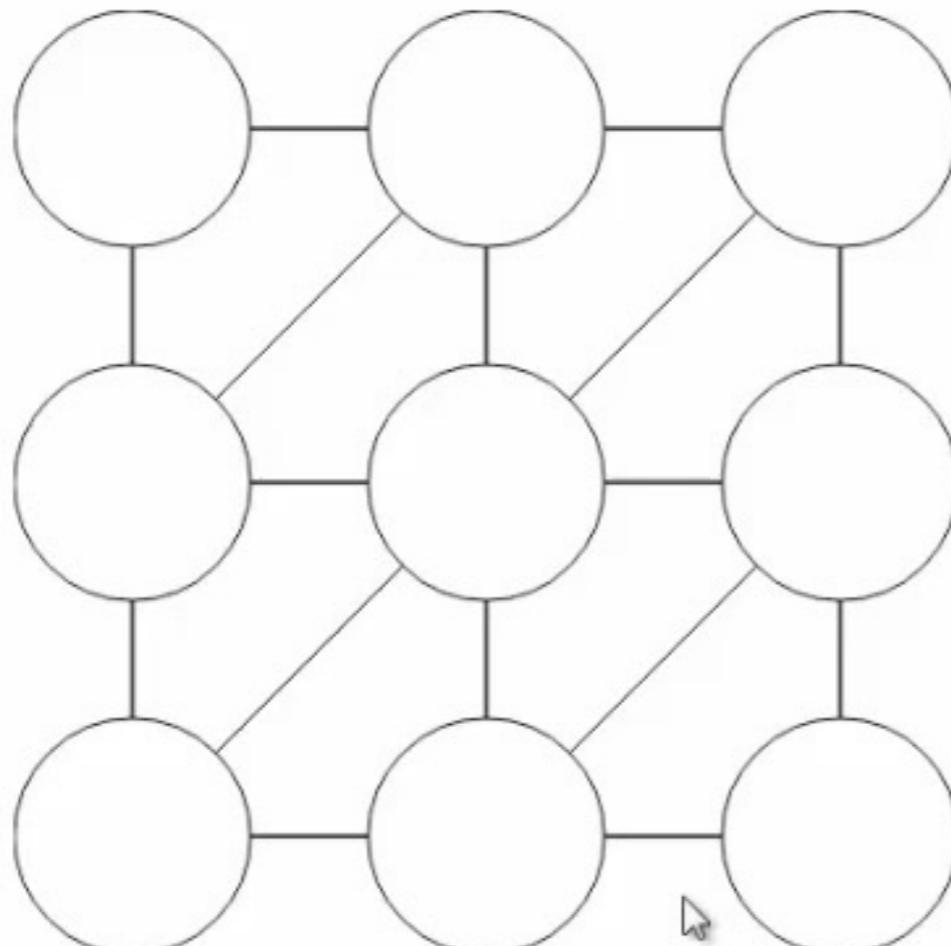


Backtracking Search

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
    return RECURSIVE-BACKTRACKING({ }, csp)
function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
    if assignment is complete then return assignment
    var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
    for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
        if value is consistent with assignment given CONSTRAINTS[csp] then
            add  $\{var = value\}$  to assignment
            result  $\leftarrow$  RECURSIVE-BACKTRACKING(assignment, csp)
            if result  $\neq$  failure then return result
            remove  $\{var = value\}$  from assignment
    return failure
```

- ❖ Backtracking = DFS + variable-ordering + fail-on-violation
- ❖ What are the choice points?

Video of Demo Coloring – Backtracking



Reset Prev Pause Next Play Faster

Graph

Simple

Algorithm

Backtracking

Ordering

- None
- MRV
- MRV with LCV

Filtering

- None
- Forward Checking
- Arc Consistency

Speed

Speedup

1 x

Frame Delay

700

Improving Backtracking

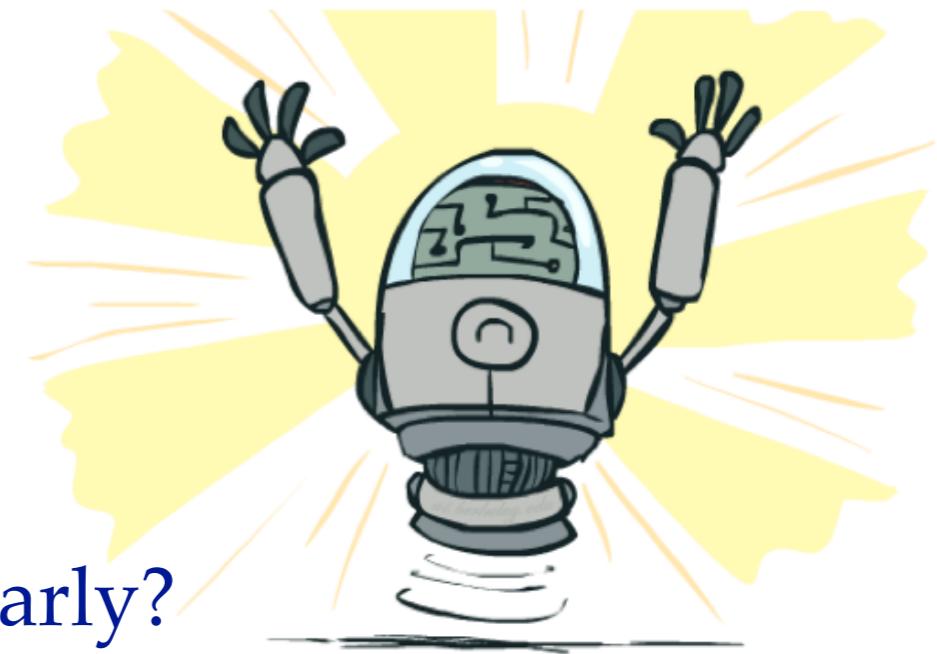
- ❖ General-purpose ideas give huge gains in speed

- ❖ Ordering:

- ❖ Which variable should be assigned next?
 - ❖ In what order should its values be tried?

- ❖ Filtering: Can we detect inevitable failure early?

- ❖ Structure: Can we exploit the problem structure?

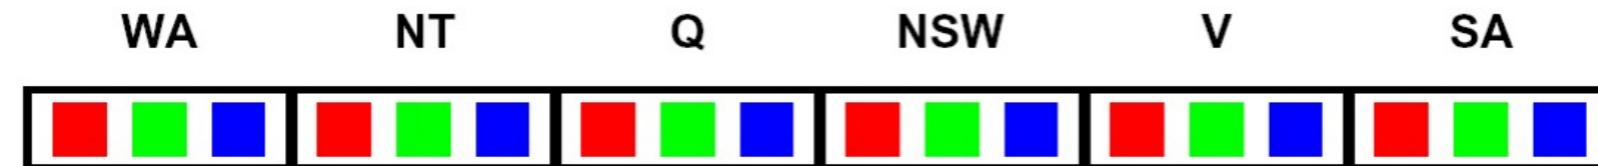


Filtering

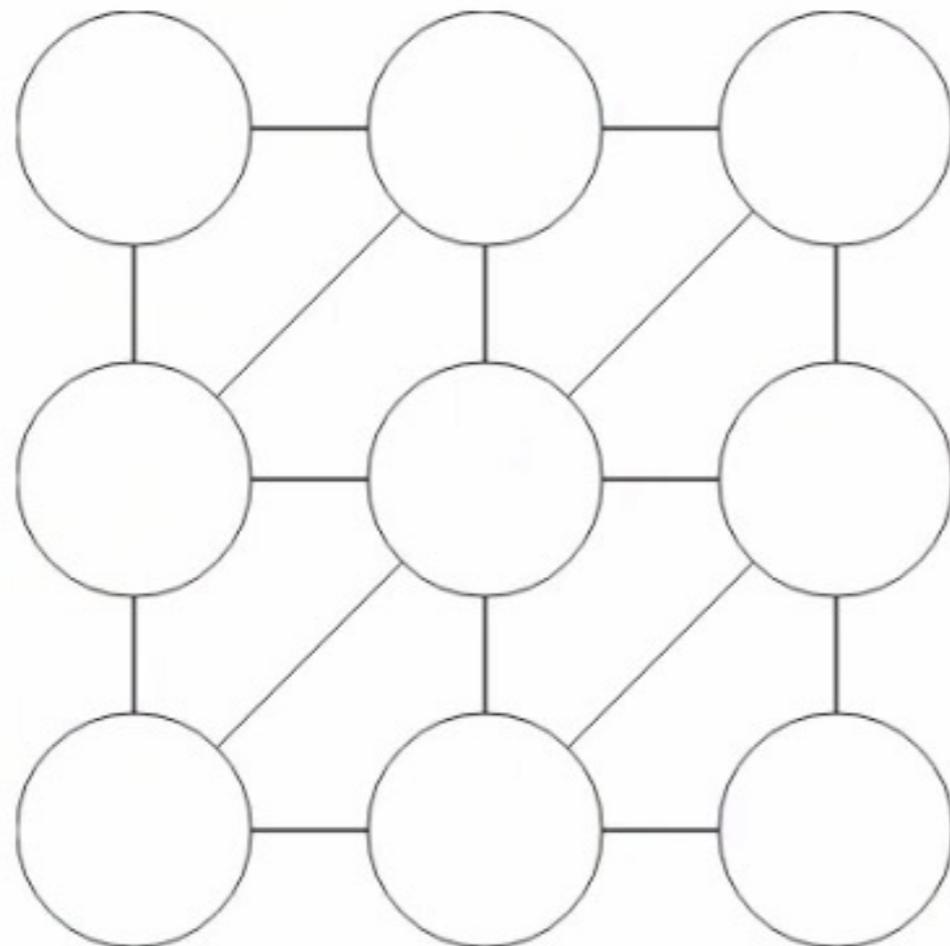


Filtering: Forward Checking

- ❖ Filtering: Keep track of domains for unassigned variables and cross off bad options
- ❖ Forward checking: Cross off values that violate a constraint when added to the existing assignment



Video of Demo Coloring – Backtracking with Forward Checking



Reset Prev Pause Next Play Faster

Graph

Simple ▾

Algorithm

Backtracking ▾

Ordering

- None
- MRV
- MRV with LCV

Filtering

- None
- Forward Checking
- Arc Consistency

Speed

Speedup Frame Delay
1 X 700

Filtering: Constraint Propagation

- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:

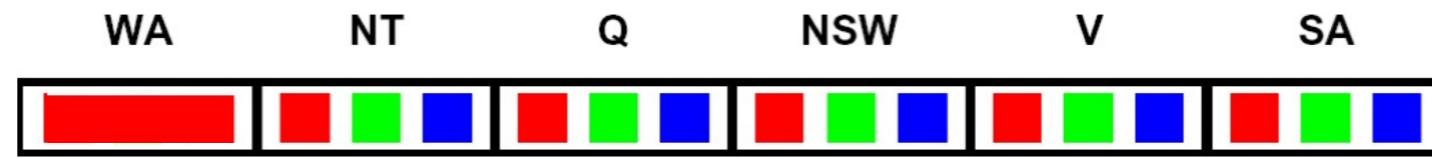


WA	NT	Q	NSW	V	SA
Red	Green	Blue	Red	Green	Blue
Red		Green	Blue	Red	Green
Red		Blue	Green	Red	Blue

- NT and SA cannot both be blue!
- Why didn't we detect this yet?
- Constraint propagation*: reason from constraint to constraint

Consistency of A Single Arc

- An arc $X \rightarrow Y$ is **consistent** iff for *every* x in the tail there is *some* y in the head which could be assigned without violating a constraint

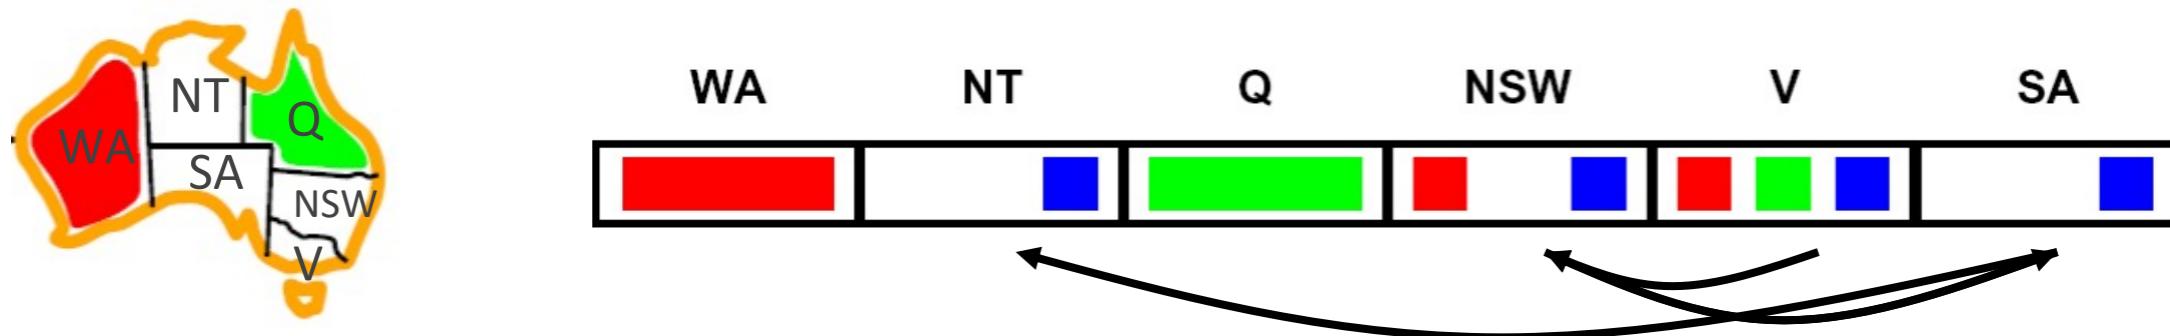


Delete from the tail!

- Forward checking: Enforcing consistency of arcs pointing to each new assignment

Arc Consistency of an Entire CSP

- ❖ A simple form of propagation makes sure **all** arcs are consistent:



- ❖ Important: If X loses a value, neighbors of X need to be rechecked!
- ❖ Arc consistency detects failure earlier than forward checking
- ❖ Can be run as a preprocessor or after each assignment
- ❖ What's the downside of enforcing arc consistency?

*Remember:
Delete from
the tail!*

Enforcing Arc Consistency in a CSP

```
function AC-3( csp) returns the CSP, possibly with reduced domains
  inputs: csp, a binary CSP with variables  $\{X_1, X_2, \dots, X_n\}$ 
  local variables: queue, a queue of arcs, initially all the arcs in csp

  while queue is not empty do
     $(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\text{queue})$ 
    if REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) then
      for each  $X_k$  in NEIGHBORS[ $X_i$ ] do
        add  $(X_k, X_i)$  to queue

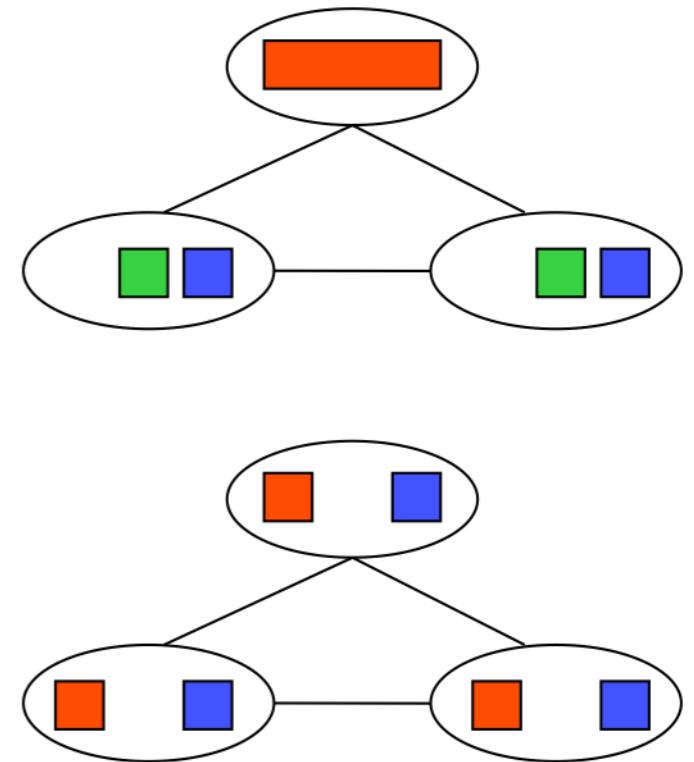
  function REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) returns true iff succeeds
    removed  $\leftarrow \text{false}$ 
    for each  $x$  in DOMAIN[ $X_i$ ] do
      if no value  $y$  in DOMAIN[ $X_j$ ] allows  $(x, y)$  to satisfy the constraint  $X_i \leftrightarrow X_j$ 
        then delete  $x$  from DOMAIN[ $X_i$ ]; removed  $\leftarrow \text{true}$ 
    return removed
```

- ❖ Runtime: $O(n^2d^3)$, can be reduced to $O(n^2d^2)$
- ❖ ... but detecting all possible future problems is NP-hard – why?

Limitations of Arc Consistency

❖ After enforcing arc consistency:

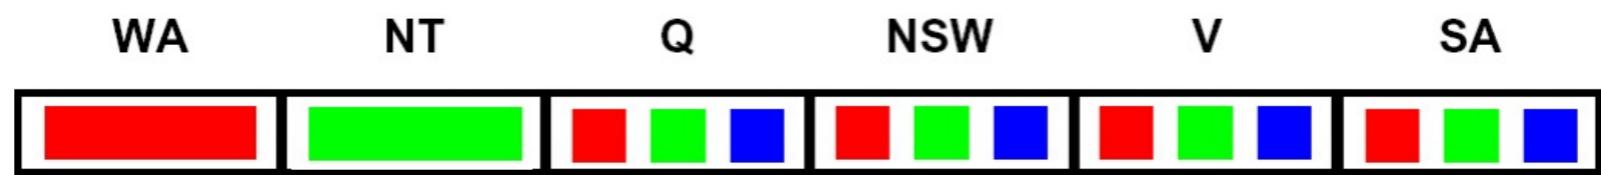
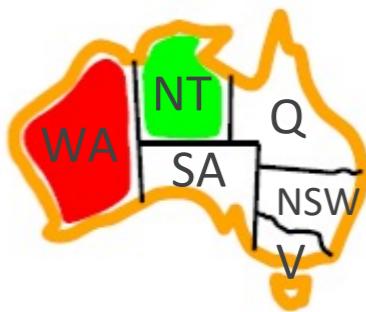
- ❖ Can have one solution left
 - ❖ Can have multiple solutions left
 - ❖ Can have no solutions left (and not know it)
- ❖ Arc consistency still runs inside a backtracking search!



What went wrong here?

Quiz: What is the color of NSW?

- ❖ Run arc consistency on this example:



- ❖ Choose a color
 - ❖ Red
 - ❖ Green
 - ❖ Blue

Video of Demo Coloring – Backtracking with Forward Checking – Complex Graph

<img alt="A complex graph visualization for a coloring demonstration. The graph consists of 20 nodes arranged in four columns of five nodes each. Each node is a circle containing three colored dots (blue, red, green). Edges connect nodes between adjacent columns and between nodes in the same column. A vertical edge connects the first node of the first column to the last node of the fourth column. On the left side, there are six buttons: 'Reset', 'Prev', 'Pause' (with a mouse cursor pointing to it), 'Next', 'Play', and 'Faster'. To the right of the graph are several configuration options: 'Graph' set to 'Complex', 'Algorithm' set to 'Backtracking', 'Ordering' with radio buttons for 'None', 'MRV', and 'MRV with LCV' (the latter is selected), 'Filtering' with radio buttons for 'None', 'Forward Checking' (selected), and 'Arc Consistency', and 'Speed' controls for 'Speedup' (set to 1) and 'Frame Delay' (set to 700).</div>

Video of Demo Coloring – Backtracking with Arc Consistency – Complex Graph

A complex graph for a 3-coloring problem. The graph consists of 18 nodes arranged in three columns of six nodes each. Each node is a circle containing three smaller circles, representing a state with three possible values (blue, red, green). The graph has the following connections:

- The first column has a single vertical connection from the top node to the bottom node.
- The second column has two diagonal connections from the top-left node to the bottom-right node.
- The third column has two diagonal connections from the top-left node to the bottom-right node.
- Horizontal connections exist between adjacent nodes in each row.
- Vertical connections exist between corresponding nodes in adjacent columns (e.g., top node of the first column connects to the top node of the second column).

Graph: Complex

Algorithm: Backtracking

Ordering:

- None
- MRV
- MRV with LCV

Filtering:

- None
- Forward Checking
- Arc Consistency

Speed:

Speedup x Frame Delay

Reset Prev Pause Next Play Faster

Ordering

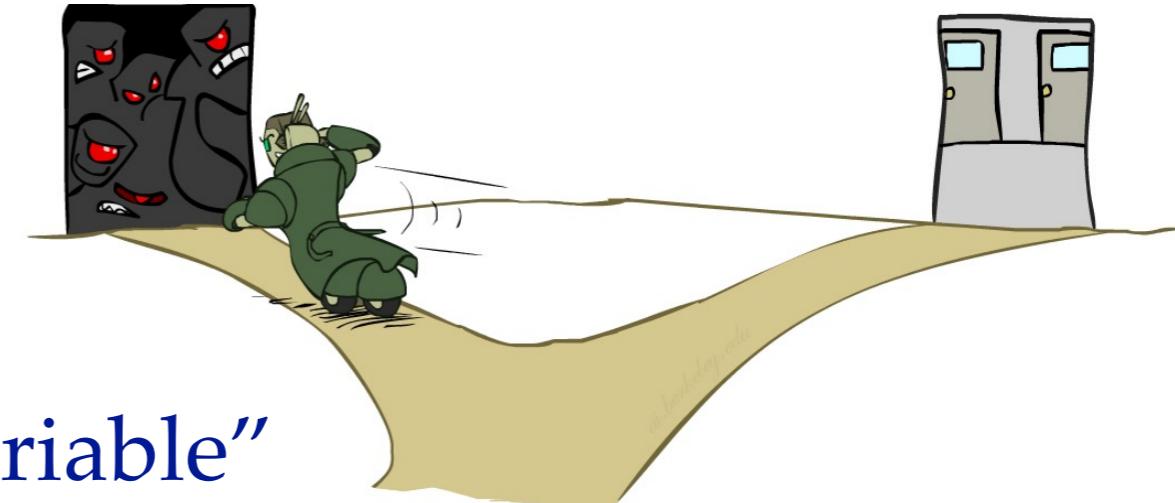


Ordering: Minimum Remaining Values

- ❖ Variable Ordering: Minimum remaining values (MRV):
 - ❖ Choose the variable with the fewest legal left values in its domain



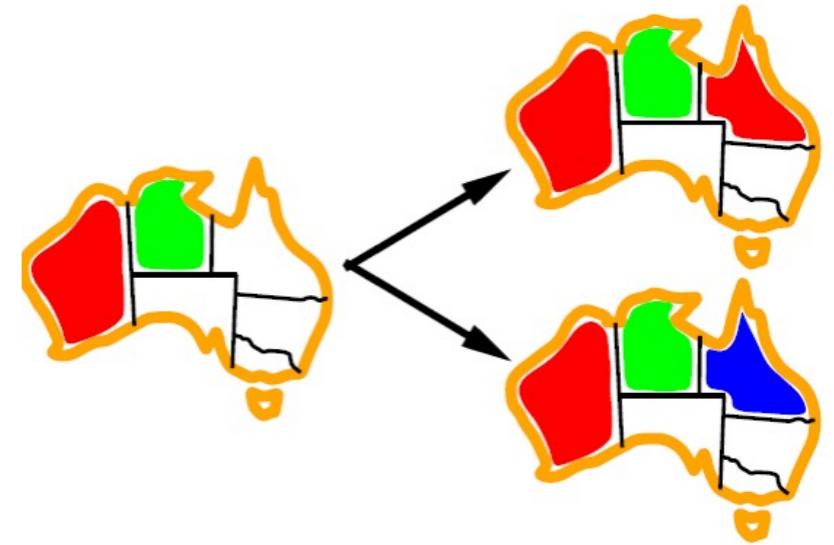
- ❖ Why min rather than max?
- ❖ Also called “most constrained variable”
- ❖ “Fail-fast” ordering



Ordering: Least Constraining Value

- ❖ Value Ordering: Least Constraining Value

- ❖ Given a choice of variable, choose the *least constraining value*
- ❖ I.e., the one that rules out the fewest values in the remaining variables
- ❖ Note that it may take some computation to determine this! (E.g., rerunning filtering)



- ❖ Why least rather than most?

- ❖ Combining these ordering ideas makes 1000 queens feasible